

Name and ID

Yunke Zhu 1330327

HW04 Code

You will complete the following notebook, as described in the PDF for Homework 04 (included in the download with the starter code). You will submit:

1. This notebook file, along with your COLLABORATORS.txt file, to the Gradescope link for code.
2. A PDF of this notebook and all of its output, once it is completed, to the Gradescope link for the PDF.

Please report any questions to the [class Piazza page \(https://piazza.com/class/k0grsypt15j73g\)](https://piazza.com/class/k0grsypt15j73g).

Import required libraries

```
In [1]: import os
import numpy as np
import pandas as pd
import time
import warnings

from sklearn.neural_network import MLPClassifier

from matplotlib import pyplot as plt
import seaborn as sns

from MLPClassifierWithSolverLBFGS import MLPClassifierLBFGS

from viz_tools_for_binary_classifier import plot_pretty_probabilities_
for_clf

%matplotlib inline
```

Load data

```
In [2]: # Load data
x_tr_N2 = np.loadtxt('./data_xor/x_train.csv', skiprows=1, delimiter=',')
x_te_N2 = np.loadtxt('./data_xor/x_test.csv', skiprows=1, delimiter=',')

y_tr_N = np.loadtxt('./data_xor/y_train.csv', skiprows=1, delimiter=',')
y_te_N = np.loadtxt('./data_xor/y_test.csv', skiprows=1, delimiter=',')

assert x_tr_N2.shape[0] == y_tr_N.shape[0]
assert x_te_N2.shape[0] == y_te_N.shape[0]
```

Problem 1: MLP size [2] with activation ReLU and L-BFGS solver

```
In [29]: # TODO edit this block to run from 16 different random_states
# Save each run's trained classifier object in a list

n_runs = 16
mlp_relu_results = []
start_time_sec = time.time()
lossCurve_LBFGS_RELU = []
for i in range(n_runs):
    mlp_lbfgs = MLPClassifierLBFGS(
        hidden_layer_sizes=[2],
        activation='relu',
        alpha=0.0001,
        max_iter=200, tol=1e-6,
        random_state=i
    )
    with warnings.catch_warnings(record=True) as warn_list:
        mlp_lbfgs.fit(x_tr_N2, y_tr_N)
    elapsed_time_sec = time.time() - start_time_sec
    print('finished LBFGS run %2d/%d after %6.1f sec | %3d iters | %s'
          | loss %.3f' % (
              i+1, n_runs, elapsed_time_sec,
              len(mlp_lbfgs.loss_curve_),
              'converged' if mlp_lbfgs.did_converge else 'NOT converged',
              mlp_lbfgs.loss_))
    lossCurve_LBFGS_RELU.append(mlp_lbfgs.loss_curve_)
    mlp_relu_results.append(mlp_lbfgs)
```


```

finished LBFGS run 1/16 after 0.0 sec | 29 iters | converged
| loss 0.347
finished LBFGS run 2/16 after 0.1 sec | 30 iters | converged
| loss 0.477
finished LBFGS run 3/16 after 0.1 sec | 21 iters | converged
| loss 0.347
finished LBFGS run 4/16 after 0.1 sec | 35 iters | converged
| loss 0.347
finished LBFGS run 5/16 after 0.1 sec | 29 iters | converged
| loss 0.347
finished LBFGS run 6/16 after 0.1 sec | 29 iters | converged
| loss 0.000
finished LBFGS run 7/16 after 0.1 sec | 23 iters | converged
| loss 0.000
finished LBFGS run 8/16 after 0.1 sec | 37 iters | converged
| loss 0.347
finished LBFGS run 9/16 after 0.2 sec | 15 iters | converged
| loss 0.347
finished LBFGS run 10/16 after 0.2 sec | 26 iters | converged
| loss 0.000
finished LBFGS run 11/16 after 0.2 sec | 36 iters | converged
| loss 0.347
finished LBFGS run 12/16 after 0.2 sec | 27 iters | converged
| loss 0.477
finished LBFGS run 13/16 after 0.2 sec | 39 iters | converged
| loss 0.000
finished LBFGS run 14/16 after 0.3 sec | 29 iters | converged
| loss 0.347
finished LBFGS run 15/16 after 0.3 sec | 25 iters | converged
| loss 0.347
finished LBFGS run 16/16 after 0.3 sec | 30 iters | converged
| loss 0.347

```

1 (a): Visualize probabilistic predictions in 2D feature space for ReLU + L-BFGS

```
In [4]: fig, ax_grid = plt.subplots(nrows=4, ncols=4, figsize=(16, 16))
grid_n_runs = [_ for _ in range(n_runs)]
grid_n_runs = np.array(grid_n_runs).reshape(4,4)
for i in range(4):
    for j in range(4):
        plot_pretty_probabilities_for_clf(mlp_relu_results[grid_n_runs
[i][j]], x_tr_N2, y_tr_N, ax=ax_grid[i,j])
```



1 (b): What fraction of runs reach 0 training error? What happens to the others? Describe how rapidly (or slowly) things seem to converge.

Answer: There are four prediction reach 0 training error, which is a fraction of $\frac{4}{16}$ meaning that 25% runs can predict absolutely right without any errors. Others runs error rate are about 25%, which means that we can predict 75% training data in these runs correctly. RELU+L-BFGS is rapidly converged. The longest iteration taken is under 37 iterations, the average times of iteration taken is 29.

Problem 2: MLP size [2] with activation Logistic and L-BFGS solver

```
In [31]: n_runs = 16
mlp_logistic_results = []
start_time_sec = time.time()
lossCurve_LBFGS_LOGISTIC = []
for i in range(n_runs):
    mlp_lbfgs = MLPClassifierLBFGS(
        hidden_layer_sizes=[2],
        activation='logistic',
        alpha=0.0001,
        max_iter=200, tol=1e-6,
        random_state=i
    )
    with warnings.catch_warnings(record=True) as warn_list:
        mlp_lbfgs.fit(x_tr_N2, y_tr_N)
    elapsed_time_sec = time.time() - start_time_sec
    print('finished LBFGS run %2d/%d after %6.1f sec | %3d iters | %s'
          | loss %.3f' % (
              i+1, n_runs, elapsed_time_sec,
              len(mlp_lbfgs.loss_curve_),
              'converged' if mlp_lbfgs.did_converge else 'NOT converged',
              mlp_lbfgs.loss_))
    lossCurve_LBFGS_LOGISTIC.append(mlp_lbfgs.loss_curve_)
    mlp_logistic_results.append(mlp_lbfgs)
```


```

finished LBFGS run 1/16 after 0.0 sec | 58 iters | converged
| loss 0.000
finished LBFGS run 2/16 after 0.1 sec | 125 iters | converged
| loss 0.347
finished LBFGS run 3/16 after 0.1 sec | 45 iters | converged
| loss 0.347
finished LBFGS run 4/16 after 0.1 sec | 76 iters | converged
| loss 0.000
finished LBFGS run 5/16 after 0.2 sec | 40 iters | converged
| loss 0.477
finished LBFGS run 6/16 after 0.2 sec | 42 iters | converged
| loss 0.000
finished LBFGS run 7/16 after 0.2 sec | 50 iters | converged
| loss 0.000
finished LBFGS run 8/16 after 0.2 sec | 14 iters | converged
| loss 0.537
finished LBFGS run 9/16 after 0.2 sec | 91 iters | converged
| loss 0.347
finished LBFGS run 10/16 after 0.3 sec | 100 iters | converged
| loss 0.347
finished LBFGS run 11/16 after 0.3 sec | 65 iters | converged
| loss 0.347
finished LBFGS run 12/16 after 0.3 sec | 104 iters | converged
| loss 0.347
finished LBFGS run 13/16 after 0.4 sec | 60 iters | converged
| loss 0.347
finished LBFGS run 14/16 after 0.4 sec | 33 iters | converged
| loss 0.478
finished LBFGS run 15/16 after 0.4 sec | 53 iters | converged
| loss 0.000
finished LBFGS run 16/16 after 0.4 sec | 61 iters | converged
| loss 0.000

```

2 (a): Visualize probabilistic predictions in 2D feature space for Logistic Sigmoid + L-BFGS

```
In [7]: fig, ax_grid = plt.subplots(nrows=4, ncols=4, figsize=(16, 16))
grid_n_runs = [_ for _ in range(n_runs)]
grid_n_runs = np.array(grid_n_runs).reshape(4,4)
for i in range(4):
    for j in range(4):
        plot_pretty_probabilities_for_clf(mlp_logistic_results[grid_n_runs[i][j]], x_tr_N2, y_tr_N, ax=ax_grid[i,j])
```



2 (b): What fraction of runs reach 0 training error? What happens to the others? Describe how rapidly (or slowly) things seem to converge.

Answer: 37.5% ($\frac{6}{16}$) of the runs reach 0 training error. As for Other runs the error rate is between 0.32 to 0.25. For most cases, these models has an error rate about 25%(a little more than 25%). This model takes more time to converge compared to RELU, the iterations taken is 14 and 125, and average iterations taken is about 63 which is more than slower than RELU model using the same L-BFGS solver.

Problem 3: MLP size [2] with activation ReLU and SGD solver


```
In [32]: # TODO edit this block to do 16 different runs (each with different random_state value)
# Save each run's trained classifier object in a list

mlp_relu_sgd_results = []
start_time_sec = time.time()
lossCurve_SGD_RELU = []
for i in range(n_runs):
    mlp_sgd = MLPClassifier(
        hidden_layer_sizes=[2],
        activation='relu',
        alpha=0.0001,
        max_iter=400, tol=1e-8,
        random_state=i,
        solver='sgd', batch_size=10,
        learning_rate='adaptive', learning_rate_init=0.1, momentum=0.0
    )
    with warnings.catch_warnings(record=True) as warn_list:
        mlp_sgd.fit(x_tr_N2, y_tr_N)
        mlp_sgd.did_converge = True if len(warn_list) == 0 else False
        elapsed_time_sec = time.time() - start_time_sec
        print('finished SGD run %2d/%d after %6.1f sec | %3d epochs | %s |'
              'loss %.3f' % (
                  i+1, n_runs, elapsed_time_sec,
                  len(mlp_sgd.loss_curve_),
                  'converged' if mlp_sgd.did_converge else 'NOT converged',
                  mlp_sgd.loss_))
        lossCurve_SGD_RELU.append(mlp_sgd.loss_curve_)
        mlp_relu_sgd_results.append(mlp_sgd)
```


finished SGD run	1/16 after	4.8 sec	267 epochs	converged
loss 0.347				
finished SGD run	2/16 after	10.2 sec	307 epochs	converged
loss 0.478				
finished SGD run	3/16 after	14.6 sec	239 epochs	converged
loss 0.347				
finished SGD run	4/16 after	21.7 sec	400 epochs	NOT converged
loss 0.001				
finished SGD run	5/16 after	26.6 sec	275 epochs	converged
loss 0.347				
finished SGD run	6/16 after	33.6 sec	400 epochs	NOT converged
loss 0.001				
finished SGD run	7/16 after	40.7 sec	400 epochs	NOT converged
loss 0.001				
finished SGD run	8/16 after	45.5 sec	273 epochs	converged
loss 0.347				
finished SGD run	9/16 after	49.5 sec	219 epochs	converged
loss 0.347				
finished SGD run	10/16 after	56.7 sec	400 epochs	NOT converged
loss 0.001				
finished SGD run	11/16 after	63.9 sec	394 epochs	converged
loss 0.478				
finished SGD run	12/16 after	70.9 sec	400 epochs	NOT converged
loss 0.478				
finished SGD run	13/16 after	77.9 sec	400 epochs	NOT converged
loss 0.002				
finished SGD run	14/16 after	83.2 sec	304 epochs	converged
loss 0.347				
finished SGD run	15/16 after	89.3 sec	331 epochs	converged
loss 0.347				
finished SGD run	16/16 after	96.3 sec	400 epochs	NOT converged
loss 0.001				

3 (a): Visualize probabilistic predictions in 2D feature space for ReLU + SGD

```
In [9]: fig, ax_grid = plt.subplots(nrows=4, ncols=4, figsize=(16, 16))
grid_n_runs = [_ for _ in range(n_runs)]
grid_n_runs = np.array(grid_n_runs).reshape(4,4)
for i in range(4):
    for j in range(4):
        plot_pretty_probabilities_for_clf(mlp_relu_sgd_results[grid_n_runs[i][j]], x_tr_N2, y_tr_N, ax=ax_grid[i,j])
```



3 (b): What fraction of runs reach 0 training error? What happens to the others? Describe how rapidly (or slowly) things seem to converge.

Answer: 37.5% ($\frac{6}{16}$) of the runs reach 0 training error. As for Other runs the error rate is all 0.25. This model takes much more time to converge compared to L-BFGS using the same activation function, there are 7 models that are not even converged after 400 iterations. And the time consuming is more than L-BFGS models since this model needs more time for converge.

3 (c): What is most noticeably different between SGD with batch size 10 and the previous L-BFGS in part 1 (using the same ReLU activation function)? Why, do you believe, these differences exist?

Answer: Compared with L-BFGS in part1, SGD has more predictions with 0 training errors (37.5% and 25%). And SGD solver takes more time running 16 models with the same size of data. And this is also the reason that SGD appears some NOT Converged runs. \ The main reason causing the difference is that SGD takes more time to converge the models. The root reason is SGD take the advantage of first derivative and L-BFGS takes both first derivative and second derivative. This helps L-BFGS converge more rapidly. The advantage of L-BFGS is that it can optimize the model by rapidly converge using it's featrue of second derivative, which SGD doesn't have.

Problem 4: MLP size [2] with activation Logistic and SGD solver

```

In [34]: # TODO edit to do 16 runs of SGD, like in previous step, but with LOGI
          # STIC activation
          # TODO edit this block to do 16 different runs (each with different ra
          # ndom_state value)
          # Save each run's trained classifier object in a list


mlp_logistic_sgd_results = []
start_time_sec = time.time()
lossCurve_SGD_LOGISTIC = []
for i in range(n_runs):
    mlp_sgd = MLPClassifier(
        hidden_layer_sizes=[2],
        activation='logistic',
        alpha=0.0001,
        max_iter=400, tol=1e-8,
        random_state=i,
        solver='sgd', batch_size=10,
        learning_rate='adaptive', learning_rate_init=0.1, momentum=0.0
    )
    with warnings.catch_warnings(record=True) as warn_list:
        mlp_sgd.fit(x_tr_N2, y_tr_N)
        mlp_sgd.did_converge = True if len(warn_list) == 0 else False
        elapsed_time_sec = time.time() - start_time_sec
        print('finished SGD run %2d/%d after %6.1f sec | %3d epochs | %s |
loss %.3f' % (
            i+1, n_runs, elapsed_time_sec,
            len(mlp_sgd.loss_curve_),
            'converged' if mlp_sgd.did_converge else 'NOT converge
d',
            mlp_sgd.loss_))
        lossCurve_SGD_LOGISTIC.append(mlp_sgd.loss_curve_)
        mlp_logistic_sgd_results.append(mlp_sgd)

```

finished SGD run	1/16 after	2.3 sec	161 epochs	converged
loss 0.693				
finished SGD run	2/16 after	7.7 sec	400 epochs	NOT converged
loss 0.005				
finished SGD run	3/16 after	13.2 sec	400 epochs	NOT converged
loss 0.005				
finished SGD run	4/16 after	16.1 sec	215 epochs	converged
loss 0.693				
finished SGD run	5/16 after	21.5 sec	400 epochs	NOT converged
loss 0.351				
finished SGD run	6/16 after	27.0 sec	400 epochs	NOT converged
loss 0.005				
finished SGD run	7/16 after	33.2 sec	400 epochs	NOT converged
loss 0.005				
finished SGD run	8/16 after	39.1 sec	400 epochs	NOT converged
loss 0.351				
finished SGD run	9/16 after	44.8 sec	400 epochs	NOT converged
loss 0.351				
finished SGD run	10/16 after	50.2 sec	400 epochs	NOT converged
loss 0.351				
finished SGD run	11/16 after	52.0 sec	124 epochs	converged
loss 0.693				
finished SGD run	12/16 after	57.4 sec	400 epochs	NOT converged
loss 0.005				
finished SGD run	13/16 after	62.9 sec	400 epochs	NOT converged
loss 0.005				
finished SGD run	14/16 after	68.4 sec	400 epochs	NOT converged
loss 0.353				
finished SGD run	15/16 after	73.9 sec	400 epochs	NOT converged
loss 0.007				
finished SGD run	16/16 after	79.3 sec	400 epochs	NOT converged
loss 0.005				

4(a): Visualize probabilistic predictions in 2D feature space for Logistic + SGD

```
In [11]: fig, ax_grid = plt.subplots(nrows=4, ncols=4, figsize=(16, 16))
grid_n_runs = [_ for _ in range(n_runs)]
grid_n_runs = np.array(grid_n_runs).reshape(4,4)
for i in range(4):
    for j in range(4):
        plot_pretty_probabilities_for_clf(mlp_logistic_sgd_results[grid_n_runs[i][j]], x_tr_N2, y_tr_N, ax=ax_grid[i,j])
```



4 (b): What fraction of runs reach 0 training error? What happens to the others? Describe how rapidly (or slowly) things seem to converge.

Answer: $50\%(\frac{8}{16})$ As for Other runs the error rate is between 0.498 and 0.371, which means that for those models exists training error, there is at least 37% error rate. This model takes much more time to converge compared to L-BFGS using the same activation function(logistic sigmoid), there are only two runs that are converged. And for the rest runs, they are not converged after 400 iterations. And the time consuming is more than L-BFGS models since this model needs more time for converge.

4 (c): What is most noticeably different between SGD with batch size 10 and the previous L-BFGS runs in part 2 (using the same logistic activation function)? Why, do you believe, these differences exist?

Answer: Compared with L-BFGS slover using the same logistic sigmoid activation function, this model contains much more cases with 0 trainig errors. But this case is similar as the last one but much more time consuming , and the main difference is that solver SGD is not good at converge. Like mentioned before, SGD normally use first derivative while L-BFGS use both first and second derivative that helps it converge rapidly.

Problem 5: Comparing loss_curves

5 (a): Plot loss_curves for each method in 2 x 2 subplot grid


```
In [35]: fig, ax_grid = plt.subplots(nrows=2, ncols=2, sharex=True, sharey=True
, figsize=(8,5))

for l in lossCurve_LBFGS_RELU:
    ax_grid[0][0].plot(l)

for l in lossCurve_LBFGS_LOGISTIC:
    ax_grid[0][1].plot(l)
for l in lossCurve_SGD_RELU:
    ax_grid[1][0].plot(l)
for l in lossCurve_SGD_LOGISTIC:
    ax_grid[1][1].plot(l)

# TODO plot 16 curves for each of the 2x2 settings of solver and activ
ation
ax_grid[0,0].set_title('L-BFGS ReLU')
ax_grid[0,1].set_title('L-BFGS Logistic')

ax_grid[1,0].set_title('SGD ReLU')
ax_grid[1,1].set_title('SGD Logistic')
plt.ylim([0, 1.0]); # keep this y limit so it's easy to compare across
plots
```



5 (b): From this overview plot (plus your detailed plots from prior steps), which activation function seems easier to optimize, the ReLU or the Logistic Sigmoid? Which requires most iterations in general?

Answer: RELU seems easier to optimize since the left two models' curves(loss curves) using ReLU activation function drops rapidly than using a logistic sigmoid activation function. \ And SGD with logistic sigmoid activation function needs most iterations in general because for more loss curves in other models it needs more times of iterations in general(the length of each curve stands for the iteration it needs).

5 (c): Are you convinced that one activation function is always easier to optimize? Suggest 3 additional experimental comparisons that would be informative.

Answer: I'm not convinced that using ReLU is always easiest to optimize.\ We can build several additional experimental comparisons, for example:\

1. Using several different dataset containing different number of features.
2. Using more activation functions except ReLU and Logistic Sigmoid, such like Softmax, Tanh, and compare their performance.
3. Using more appropriate solver, such as DFP, to check if ReLU only performs well under L-BFGS and SGD.